

ReSketch: Corpus-Guided Semantic Regex Synthesis

Berk Cakar
Course Project Final Report
ECE 663: Program Synthesis
Purdue University
West Lafayette, IN, USA

Abstract

Writing a correct regular expression demands low-level syntactic knowledge that is often far removed from the user’s intent. A user may know that a field should hold an email address or a CVV code, yet translating that knowledge into a precise pattern remains error-prone. ReSketch addresses this gap with *semantic sketches*: partial regexes whose open positions are typed holes such as $\{\square : \text{email_address}\}$ or $\{\square : \text{cvv}\}$. Given a sketch and a set of positive and negative examples, ReSketch retrieves reusable regex components for each hole, decomposes the global examples into per-hole evidence, and searches over candidate components and refinements to assemble a validated regex. On a suite of nine demonstration tasks spanning single-hole, multi-hole, and ambiguous-decomposition scenarios, the prototype solves all tasks under default settings; a decomposition ablation shows that hole-local evidence is necessary for multi-hole synthesis.

1 Introduction

Regular expressions are a compact notation for describing sets of strings, widely used for validation, data extraction, log processing, and input sanitization. Yet empirical studies find that developers struggle to read, write, and maintain them [12]. The consequences go beyond readability: regex mistakes cause string-processing bugs and security vulnerabilities such as Regular Expression Denial of Service [2, 6, 8, 16]. A user may know that a field should hold an email address, an IPv4 address, or a payment-related token, but translating that semantic intent into a correct pattern remains difficult.

Program synthesis can help. Instead of writing the final expression directly, a user can supply examples, natural language, or partial specifications and let a synthesizer search for a consistent program. Programming-by-example has proved effective for string transformations in spreadsheets [9], and several systems apply the idea to regex synthesis from examples and multimodal specifications [5, 11]. Semantic regex synthesis goes one step further: practical extraction tasks often require domain-level concepts (dates, currencies, IP addresses) that plain regex syntax cannot express directly [4]. Examples alone, however, leave the intended semantic category implicit. A handful of positive and negative strings may separate one candidate regex from another, but they do not tell the synthesizer that the user wanted a credit-card field rather than an arbitrary digit sequence.

ReSketch introduces a different starting point: the *semantic sketch*. The user writes a partial regex with typed holes such as $\{\square : \text{email_address}\}$ or $\text{CVV} : \{\square : \text{cvv}\}$, fixing the surrounding structure while leaving semantically meaningful sub-expressions open. ReSketch completes those holes by retrieving reusable regex components from an annotated corpus, a language model, or

both. This design reflects observed developer practice—regexes are frequently copied and adapted [7]—and is consistent with recent evidence that both reuse and LLM-based generation are competitive strategies for regex composition [3]. Retrieval proposes candidate components; examples remain the behavioral specification that drives decomposition, ranking, and validation.

The resulting system is both a research vision and a course-project prototype. The vision is a corpus-guided synthesizer backed by an annotated corpus of semantically labeled regexes. The prototype implements the core synthesis pipeline behind a pluggable provider interface, so retrieval can draw on different component sources without changing the downstream algorithm.

This report makes four contributions:

- (1) A *semantic sketch interface* for regex synthesis that combines typed holes with ordinary regex context.
- (2) A *hybrid retrieval architecture* that treats different component sources as interchangeable providers of regex fragments.
- (3) A *prototype synthesis engine* with example decomposition, component-guided search, scoring, validation, and interactive counterexample refinement.
- (4) A *demonstration on nine tasks* showing that the prototype solves single-hole and multi-hole tasks, and that example decomposition is critical for multi-hole assembly.

2 Background and Motivation

ReSketch draws on three lines of prior work: programming by example, which shows that users can communicate intent through concrete behaviors rather than complete programs; semantic regex systems, which make domain-level concepts explicit in the synthesis problem; and empirical studies of regex reuse, which show that developers routinely adapt existing patterns instead of writing from scratch. Together, these motivate two design choices: typed sketches should guide component retrieval, and examples should remain the final behavioral specification.

2.1 Regex Synthesis from Examples

Programming by example (PBE) is a natural fit for regex authoring: users find it easier to supply strings that should or should not match than to describe the complete language of the target expression. FlashFill showed that this interaction model works for string transformations in spreadsheets—a user gives input-output examples, and the synthesizer searches a domain-specific language for a consistent program [9]. Regex synthesis applies the same idea to a different output language: given positive and negative strings, find a regex that accepts the positives and rejects the negatives.

Early systems explored different search strategies for this problem. Bartoli et al. used genetic programming to evolve regexes

from labeled extraction examples [1]. Lee et al. searched for simple regexes consistent with positive and negative examples, using over- and under-approximations to prune the candidate space [11]. Both illustrate the central tradeoff ReSketch also faces: examples provide a precise behavioral test, but they do not by themselves shrink the search space.

More recent systems add a second modality to reduce ambiguity. Multimodal regex synthesis combines natural language with examples: the language describes high-level intent, and the examples rule out incorrect completions [5]. Sketch-driven regex generation takes this further by mapping natural language to an incomplete regex sketch that a synthesizer then fills using examples [17]. ReSketch adopts this sketch-and-examples structure but changes the source of the sketch: instead of deriving it from natural language, the user names semantic holes directly, and those names drive component retrieval.

2.2 Semantic Regular Expressions

Many extraction tasks are only partly syntactic. A date or IPv4 address has a recognizable surface form, but categories like *business entity* or *city in Indiana* require knowledge about what the matched text denotes, not just how it looks. Chen et al. formalize this observation by extending regexes with semantic matching constructs of the form $\{v : \tau \mid \phi\}$, where τ names a semantic type and ϕ optionally refines it with a predicate [4]. Their Smore system synthesizes such semantic regexes from positive and negative examples via neural sketch generation, decomposition, and type-directed search.

ReSketch shares the high-level insight but makes a different engineering choice. Instead of extending the *runtime* regex language with semantic predicates, ReSketch treats semantic names as *compile-time* guidance for retrieving ordinary regex components. A hole such as $\{\square : \text{email_address}\}$ does not ask the matcher to query a semantic oracle at match time; it asks the synthesizer to retrieve candidate regexes whose intended meaning matches the label, then validate the completed expression against examples.

2.3 Why Reuse Matters

Regex authoring in practice is seldom a blank-page problem. Developers routinely search for, copy, and adapt existing expressions [7], because many validation and extraction tasks recur across applications. Reuse provides strong candidate structure, but it is not a correctness argument by itself.

ReSketch turns this observation into a controlled retrieval-and-validation strategy. Retrieved regexes are treated as hypotheses—they may come from an annotated corpus, a language model, or a combination—and must survive decomposition, search, and validation before they are accepted. Recent work suggests that both reuse and LLM-based generation are competitive approaches to regex composition [3], which motivates a provider-agnostic retrieval interface: the downstream synthesis algorithm is the same regardless of where candidates originate.

3 System Overview

ReSketch separates retrieval from validation. Semantic labels guide retrieval toward plausible regex components; examples decide which candidates are behaviorally correct. Figure 1 shows the main

workflow, from sketch parsing through decomposition, retrieval, search, and assembly to a validated regex with a provenance trace. Different retrieval sources share the same candidate-set interface, so the synthesis engine ranks, refines, and validates components uniformly regardless of their origin.

3.1 Inputs and Outputs

The user supplies two things: a *semantic sketch*—ordinary regex context with one or more typed holes such as $\{\square : \text{email_address}\}$, $\{\square : \text{cvv}\}$, or $\{\square : \text{ipv4_address}\}$ —and a set of positive and negative examples for the complete regex. Optional hole-level examples can be added when the sketch context does not uniquely split examples across holes.

The system returns a validated regex together with a machine-readable *trace*: retrieved candidates and their provenance, inferred per-hole evidence, search and pruning statistics, score breakdowns, and final validation results.

Listing 1: Example ReSketch invocation.

```
resketch synthesize \
--sketch 'CVV:{\square: cvv}' \
--pos 'CVV:123' \
--pos 'CVV:1234' \
--neg 'CVV:12' \
--neg 'CVV:abcd'
```

3.2 Pipeline

Decomposition. The pipeline starts by parsing the sketch into a regex AST with named holes, preserving fixed context as structure rather than flattening it into text. ReSketch then matches this skeleton against the global examples to infer per-hole evidence. Unique captures become *hard* local positives or negatives; ambiguous captures become *soft* evidence that influences ranking without pruning; and some negatives produce *tuple-level* constraints that span multiple holes. Each hole’s active provider then receives the semantic type, the full sketch, and the decomposed evidence, and returns a finite candidate set.

Search and assembly. Each hole is searched locally before global assembly. Retrieved regexes are parsed into an internal AST when possible, then expanded through four strategies: direct reuse, component generalization, fragment recombination, and parameterized refinement. Candidates that violate hard local evidence are discarded; the rest are ranked by semantic fit, source confidence, syntactic complexity, and strategy cost. ReSketch then runs bounded beam assembly over hole assignments, pruning partial assignments that fail global positives and complete assignments that match global negatives. The surviving regex is evaluated against all original examples. In interactive mode, the user can reject a candidate and supply counterexamples; the pipeline reruns with the strengthened specification.

3.3 Retrieval Modes

ReSketch treats retrieval as a pluggable provider interface. In *corpus mode*, the provider returns regex components whose labels, descriptions, or examples match the requested semantic type. In *LLM mode*, a language model receives the semantic type, global examples, and inferred hole evidence as prompt context and proposes candidate




Figure 1: Overview of ReSketch: examples decompose into hole evidence, candidate sources propose components, and search plus assembly produce a validated regex and trace.

components. In *hybrid mode*, both sources contribute to a single candidate set before search begins. This design lets the tool operate even when one source has limited coverage for a given category. In all modes, retrieved components are only hypotheses; they must survive local evidence checks, global assembly, and final validation before they are accepted.

4 Approach

The key design principle behind ReSketch is to keep semantic retrieval and behavioral validation separate: retrieval proposes candidate components for typed holes, but the examples and the surrounding regex context determine which components survive. Algorithm 1 gives the top-level procedure; the subsections below elaborate each step.

4.1 Semantic Sketches

A semantic sketch is a partial regex that mixes ordinary regex structure with typed holes. Formally, we write a hole as $h_i : \tau_i$, where h_i is a fresh hole identifier and τ_i is an open-vocabulary semantic label such as `email_address`, `cvv`, or `ipv4_address`. A sketch therefore has the shape of a regular expression with selected subexpressions left unspecified:

$$S ::= r \mid S_1 S_2 \mid (S_1 \mid S_2) \mid S^{m,n} \mid \{\square : \tau\}.$$

Here r ranges over concrete regex fragments supported by the implementation, including literals, character classes, regex categories such as `\d`, anchors, alternation, and repetition.

Given a sketch S , examples E^+ and E^- , and a retrieval provider P , synthesis searches for an assignment A from each hole h_i to a regex r_i . Rendering the sketch with this assignment yields $S[A]$. A solution must satisfy the global examples:

$$\forall e \in E^+. S[A](e) \quad \text{and} \quad \forall e \in E^-. \neg S[A](e).$$

The semantic type τ_i does not change regex matching semantics at run time; it constrains retrieval at synthesis time by asking the provider for a finite candidate set $C_i = P(S, h_i, \tau_i, E^+, E^-, L_i)$, where L_i summarizes any hole-local evidence inferred from decomposition. This follows the sketching tradition of leaving selected program fragments open while preserving the programmer’s surrounding

Algorithm 1 SYNTHESIZE(S, E^+, E^-)

Require: Sketch S with holes $h_1:\tau_1, \dots, h_k:\tau_k$; positive examples E^+ ; negative examples E^-

Ensure: Validated regex $S[A]$ or failure

```

1:  $\mathcal{H} \leftarrow \text{PARSESKETCH}(S)$   $\triangleright$  extract holes and regex context
2:  $\langle L^+, \tilde{L}^+, L^-, C_{\text{tn}}, C_{\text{rep}} \rangle \leftarrow \text{DECOMPOSE}(S, \mathcal{H}, E^+, E^-)$   $\triangleright$  local
   evidence and global constraints
3: for each hole  $h_i \in \mathcal{H}$  do
4:    $L_i \leftarrow (L_i^+, \tilde{L}_i^+, L_i^-)$   $\triangleright$  evidence for this hole
5:    $C_i \leftarrow \text{RETRIEVE}(h_i, \tau_i, S, E^+, E^-, L_i)$   $\triangleright$  candidate
   components
6:    $O_i \leftarrow \emptyset$ 
7:   for strategy  $\sigma \in \langle \text{REUSE}, \text{GENERALIZE}, \text{RECOMBINE}, \text{REFINE} \rangle$ 
   do
8:      $O_i \leftarrow O_i \cup \text{BOUNDEDSEARCH}(\sigma)(C_i, L_i, C_{\text{tn}})$ 
9:   end for
10:  discard options in  $O_i$  that fail hard evidence  $L_i^+$  or  $L_i^-$ 
11:  rank each  $o \in O_i$  by  $\text{fit}(o) + \lambda_{\rho} \rho(o) - \text{cost}_{\text{ast}}(o) - \text{cost}_{\text{len}}(o) -$ 
    $\text{penalty}_{\sigma}(o)$ 
12: end for
13:  $\mathcal{B} \leftarrow \text{BEAMASSEMBLE}(\{O_i\}, C_{\text{tn}}, C_{\text{rep}}, E^+, E^-)$   $\triangleright$  bounded beam
   over assignments
14: for each assignment  $A \in \mathcal{B}$  do
15:   if  $\forall e \in E^+. S[A](e)$  and  $\forall e \in E^-. \neg S[A](e)$  then
16:     return  $S[A]$  with provenance trace
17:   end if
18: end for
19: return failure

```

structure [14, 15]. In ReSketch, that surrounding structure is itself valuable evidence: fixed prefixes, separators, labels, and repeated contexts help decompose global examples into constraints on individual holes.

4.2 Hybrid Candidate Retrieval

For each hole $h_i : \tau_i$, ReSketch constructs a retrieval request containing the semantic type, the full sketch, the global examples, and any hole-local evidence inferred from decomposition. A provider returns a bounded set of components

$$C_i = \{(r, d, \rho, W_r^+, W_r^-, s)\},$$

where r is a regex fragment, d is a short description, ρ is an optional confidence score, W_r^+ and W_r^- are positive and negative witness strings supplied with the component, and s records provenance. The notation E^+ and E^- is reserved for the user’s global examples. ReSketch uses this common schema for all retrieval sources, so the later synthesis stages do not depend on whether a component came from an annotated corpus, an LLM, or both.

Corpus retrieval. returns components that have already been labeled with semantic intent. A request for `cvv` or `zip_code` retrieves components whose annotations, descriptions, or examples match that category.

LLM retrieval. treats a language model as an open-vocabulary component proposer: it receives the same request context and returns candidate regexes with descriptions, confidence values, and example witnesses. This follows a broader synthesis pattern in which informal signals propose likely program fragments while examples supply behavioral constraints [5, 13].

Hybrid retrieval. unions the candidate sets from the enabled sources, deduplicates equivalent regex strings, and preserves source provenance. The goal is not to trust more components but to give search a diverse frontier: corpus retrieval can supply stable reusable idioms, while LLM retrieval can cover labels that are sparse or absent from the corpus [3]. Regardless of origin, all candidates pass through the same local checks, transformations, and global validation.

4.3 Example Decomposition

Global examples describe the whole regex, but hole search benefits from local evidence. ReSketch derives this evidence by interpreting the sketch as a partial matcher: concrete regex structure is matched normally, while each hole captures a bounded substring. For a positive example $e \in E^+$, the matcher enumerates capture assignments

$$\mathcal{A}(e) = \{\alpha \mid S[\alpha] \text{ is consistent with } e\}.$$

Hard and soft evidence. If all assignments agree that hole h_i captured value v , ReSketch records v as *hard* positive evidence L_i^+ . For example, `CVV: 123` decomposes uniquely against `CVV: {\square: cvv}`, yielding 123 as a local witness. When several splits are possible—adjacent holes or ambiguous alternations—the captures are recorded as *soft* positives $\tilde{L}_i^+$. Soft evidence improves ranking but does not prune candidates on its own.

Negative decomposition. Negatives are decomposed conservatively. A unique single-hole capture becomes hard negative evidence L_i^- . A unique multi-hole capture becomes a *tuple-negative constraint*: each value may be reasonable alone, but the joint assignment must be rejected during global assembly. Repeated occurrences of the same hole produce consistency constraints over the

captured values. Decomposition thus converts the sketch’s fixed regex context into pruning information without requiring manual hole-level labels.

4.4 Component-Guided Search

For each hole, ReSketch turns retrieved components into a local search frontier. Components that are syntactically invalid or too large are discarded; the rest are parsed into an internal regex AST so that the synthesizer can manipulate regex structure rather than raw strings. Search then produces a ranked option set O_i for hole h_i , using four strategies in increasing order of flexibility.

Direct reuse keeps a retrieved component unchanged. This is the cheapest strategy and captures the case where a provider already proposed the right regex. *Component generalization* combines pairs of retrieved components by anti-unification: for instance, compatible character classes can be widened, literal positions can be generalized to classes, and nearby repetitions can be merged into a broader quantifier. *Fragment recombination* extracts subexpressions from retrieved components and hole examples, then enumerates small expressions built with concatenation, alternation, optionality, and bounded repetition. Finally, *parameterized refinement* makes local edits to promising components, such as adjusting repeat bounds, swapping character classes, toggling anchors, or inserting optional separators suggested by the examples.

Every generated candidate is normalized, rendered back to a regex, and checked against the hard local evidence L_i^+ and L_i^- . Candidates that fail hard evidence are pruned. Soft positives $\tilde{L}_i^+$ and tuple-negative probes affect the score, but do not by themselves prove correctness. To keep the search finite, ReSketch bounds AST size, generated candidates, behavior classes, and per-hole time. It also deduplicates both syntactically identical regexes and candidates that behave the same on the observed local examples. The output of this phase is therefore not a single regex per hole, but a compact set of diverse, evidence-compatible options for global assembly.

4.5 Ranking, Validation, and CEGIS

Each hole option receives a score that balances behavioral evidence and simplicity:

$$\text{score}(o) = \text{fit}(o) + \lambda_\rho \rho(o) - \text{cost}_{\text{ast}}(o) - \text{cost}_{\text{len}}(o) - \text{penalty}_{\text{strategy}}(o).$$

The fit term rewards hard local matches, hard local rejections, soft positive matches, and tuple-negative probe rejections. Source confidence $\rho(o)$ lets retrieval sources express uncertainty, while syntactic cost and strategy penalties prefer smaller regexes and more direct explanations when examples do not distinguish alternatives.

Global assembly. Assembly searches over assignments A that choose one option from each O_i . ReSketch orders holes by available evidence and constraint participation, then runs bounded beam search over partial assignments. Tuple-negative and repeated-hole constraints prune assignments before a full regex is rendered. When automata-backed pruning applies, partial assignments are checked for positive-example feasibility and complete assignments for negative-example rejection; otherwise the engine falls back to regex-based validation. Surviving assignments are rendered as $S[A]$ and evaluated against the original E^+ and E^- .

Interactive refinement (CEGIS). In interactive mode, each synthesis round ends with a user decision: accept the candidate regex or reject it by supplying new positive or negative counterexamples. Counterexamples are appended to E^+ or E^- , and the full pipeline reruns under the strengthened specification. This follows the standard CEGIS pattern of alternating synthesis with concrete counterexamples to refine an incomplete specification [10].

5 Prototype Implementation

The prototype is a Python 3.11 command-line tool built around the provider interface and synthesis engine described above. It is a working demonstration of the full ReSketch pipeline.

5.1 Command-Line Tool

The tool exposes a CLI. The main command, `resketch synthesize`, accepts a sketch, positive and negative examples, a retrieval mode, and synthesis controls (beam size, decomposition mode, automata pruning). Auxiliary commands validate an existing regex, inspect parsed holes, print the resolved configuration, and run a small task suite.

5.2 Synthesis Engine

The engine parses sketches into an AST with explicit hole nodes, rejects unsupported non-regular constructs, and uses Python regex semantics for candidate validation. It handles embedding details such as stripping outer anchors when a retrieved component is inserted inside a larger sketch. Automata-backed pruning is implemented for the supported fullmatch subset and fails open to regex-based checks when a construct falls outside that subset.

Auditability. Every synthesis result carries enough metadata to make decisions auditable: the trace records retrieved candidates with provenance, inferred per-hole evidence, search and pruning statistics, score breakdowns, and validation outcomes. The interactive runner wraps the synthesizer in the CEGIS loop described in Section 4: a rejected candidate can be refined by adding counterexamples, after which the pipeline reruns.

5.3 Retrieval Provider Boundary

The current prototype separates retrieval from synthesis through a common candidate-provider schema. This is the mechanism that lets ReSketch treat different component sources as interchangeable: each provider receives the same request context and returns candidates in the same format. The downstream synthesis algorithm does not depend on which provider produced a candidate; once candidates are returned, all sources are normalized, searched, ranked, and validated through the same pipeline. Details about collecting and characterizing reusable regex components are outside this report’s implementation focus and are deferred to the reuse study [3].

6 Demonstration

We tested the prototype on nine synthesis tasks (Table 1) covering single-hole reuse, labeled-context decomposition, multi-hole beam assembly, ambiguous adjacent holes, repeated holes under

quantifiers, and tuple-negative constraints. Table 2 summarizes the results: under this configuration, ReSketch solves all nine tasks.

6.1 Task Suite

The suite spans three complexity levels. Three *single-hole* tasks (email-address, signed-integer, cvv-after-label) test direct reuse of retrieved components. Two *four-hole* tasks (payment-receipt, server-log) test beam assembly: the payment receipt assembles credit-card, expiration-date, CVV, and currency-amount components into one regex, while the server log requires a fragment-recombination step for the HTTP method hole because neither positive method alone covers both examples. Four remaining tasks target specific engine behaviors: ambiguous adjacent holes, a repeated hole under a quantifier, a tuple-negative constraint, and component generalization from concrete examples.

Representative synthesized regexes include `CVV:\d{3,4}` for the labeled CVV task, `[A-Z]{1,2}\d+` for the ambiguous split, and `20\d{2}-\d+` for the tuple-negative task, where the year component narrows from a four-digit pattern to a 2000s-range pattern to reject the negative `3026-42`.

6.2 Metrics

The “Cands” column in Table 2 reports the total number of local candidates evaluated across all holes. Even direct-reuse outcomes can evaluate hundreds of candidates because each retrieved component seeds bounded generalization, recombination, and refinement. Tasks with one or two holes explore 600–1 800 candidates; four-hole tasks explore 3 600–3 900 because each hole contributes its own search frontier. “Len” shows the character length of the synthesized regex, ranging from 11 for simple single-hole tasks to 86 for the four-hole server log. Decomposition achieves 100% hard-evidence coverage on all tasks except ambiguous-split, whose two adjacent holes produce only soft evidence: the trace records 16 soft positives and 2 ambiguity groups, showing that the engine distinguishes unique from ambiguous captures.

6.3 Walkthrough: Request-Log Task

To show the pipeline stages concretely, we trace a five-hole task whose sketch mixes literal regex context with typed holes:

```
\[{\[ iso_date\]}T{\[ time\]}
  req={\[ alpha\]}{\[ integer\]}
  status={\[ status_code\]}
```

Two positive examples and two negatives are supplied:

```
pos: [2026-04-27T14:05:33] req=AB123 status=200
pos: [2026-01-15T08:30:00] req=CD456 status=201
neg: [2026-04-27T14:05:33] req=123AB status=200
neg: [2026-04-27T14:05:33] req=AB123 status=999
```

Decomposition. ReSketch matches the sketch skeleton against each positive example and extracts per-hole captures. The brackets and literal labels `T`, `req=`, and `status=` act as unambiguous separators, so three holes receive hard positive evidence: `iso_date` captures `2026-04-27` and `2026-01-15`, `time` captures `14:05:33` and `08:30:00`, and `status_code` captures `200` and `201`. The two adjacent holes `alpha` and `integer` have no separator, so the boundary is ambiguous: the string `AB123` can be split as `(A,B123)`, `(AB,123)`,

Table 1: Demonstration task specifications. Each task has 1–3 positive and 1–3 negative examples; one of each is shown. Multi-hole sketches are abbreviated for readability.

Task	#E ⁺ /#E [−]	Sketch	Positive example	Negative example
email-address	2/2	{ $\square$:email_address}	john.doe@example.com	missing-at.example.com
signed-integer	3/2	{ $\square$:integer}	~42	3.14
cvv-after-label	2/2	CVV: { $\square$:cvv}	CVV:123	CVV:12
payment-receipt	2/2	CARD: { $\square$:credit_card} EXP: { $\square$:exp_date} ...	CARD:4111 . . . EXP:04/26 CVV:123 AMT:\$19.95	. . . EXP:13/26 . . .
server-log	2/2	{ $\square$:iso_date} { $\square$:ipv4} "{ $\square$:method} { $\square$:path}"	2026-04-27 192.168.0.1 "GET /api/v1/orders"	. . . "BLAH /api/v1/orders"
product-code	2/3	{ $\square$:product_code}	A12, AB12	A13, B12
ambiguous-split	2/3	{ $\square$:alpha}{ $\square$:integer}	AB123	123AB
nested-repeat	2/3	IDS: (? : { $\square$:integer}) - { $\square$:integer}	IDS:12-34-56	IDS:12-34
tuple-negative	1/1	{ $\square$:year} - { $\square$:integer}	2026-42	3026-42

Table 2: Demonstration results on nine tasks. All tasks solved successfully in this run. #H = number of typed holes; Len = synthesized regex length in characters. Strategies: DR = direct reuse, FR = fragment recombination, PR = parameterized refinement.

Task	#H	Cands	Len	Strategies
email-address	1	1 148	50	DR
signed-integer	1	808	11	DR
cvv-after-label	1	703	11	DR
payment-receipt	4	3 600	82	DR, PR
server-log	4	3 878	86	DR, FR
product-code	1	615	11	DR
ambiguous-split	2	1 660	13	DR
nested-repeat	1	1 749	19	DR
tuple-negative	2	1 636	12	DR

(AB1,23), and so on. ReSketch records all possible splits as soft positive evidence and marks two ambiguity groups.

Retrieval and search. In this run, candidate retrieval returns 5–8 regex components per hole from the active provider. Search explores all four strategies: direct reuse tries each retrieved component unchanged; fragment recombination builds small expressions from sub-fragments; parameterized refinement adjusts quantifier bounds and character classes; and component generalization attempts anti-unification of component pairs. The three hard-evidence holes prune most generated candidates quickly: for instance, `iso_date` generates 1 194 candidates, of which 1 158 are discarded because they fail the hard local examples. The two soft-evidence holes (`alpha` and `integer`) have no hard constraints to prune against, so generated candidates are not pruned by hard local evidence and are instead ranked by soft-evidence fit and syntactic cost.

Scoring and assembly. For each hole, surviving candidates are scored by semantic fit, source confidence, AST complexity cost, and strategy penalty. For the `status_code` hole, a fragment `\d*` (score 5.7) scores slightly above a retrieved component `[1-5][0-9]{2}` (score 5.5) in per-hole ranking, but during beam assembly the broader fragment is eliminated because it accepts the negative status 999. The range-restricted component survives global validation and is selected. Beam assembly combines per-hole options into global assignments, checking each complete assignment against all four examples. The validated output is:

```

\[[0-9]{4}\-[0-1][0-9]\-[0-3][0-9]\d{2}:\d{2}:\d{2}\]
req=\w*\d+ status=[1-5][0-9]{2}

```

The provenance trace records that three holes were filled by direct reuse with anchor stripping, while the `alpha` hole was filled

by fragment recombination since no single retrieved component matched the ambiguous soft evidence.

6.4 Decomposition Ablation

To measure how much example decomposition contributes, we reran four tasks under three decomposition modes: the default hard-and-soft mode, hard-only mode (soft evidence is not used for ranking), and decomposition disabled entirely. Table 3 shows the results.

Without decomposition, three of four tasks fail outright. The survivor, `ambiguous-split`, produces a narrower regex (`[A-Z]{2}\d{3}`) that overfits to the two positive examples instead of generalizing.

Table 3: Decomposition ablation on four tasks. Each cell shows the number of candidates explored; F marks a failed synthesis.

Task	Default	No decomp	Hard-only
payment-receipt	3 600	F	3 600
server-log	3 878	F	3 878
ambiguous-split	1 660	617	619
product-code	615	F	615

Hard-only mode. does not change the outcome on tasks whose decomposition is unambiguous, but it reduces the candidate count on `ambiguous-split` from 1 660 to 619 and yields a slightly less general regex (`[A-Z]{2}\d+` instead of `[A-Z]{1,2}\d+`), because soft evidence no longer guides search toward broader candidates.

These results show that decomposition is necessary for the multi-hole tasks in this suite, and that soft evidence can improve generalization when hard decomposition is ambiguous.

7 Limitations and Future Work

Three limitations bound the claims above.

Bounded search. AST size, candidate counts, behavior classes, and per-hole time are all capped by configuration. Completeness holds only within these bounds; the engine reports its completeness status in the trace so that users can tell whether search was exhaustive or truncated.

Regex subset. The supported subset excludes lookahead, lookbehind, backreferences, and other non-regular constructs. Sketches or candidates that use these features are rejected or treated as raw

regexes only when raw-regex handling is enabled; in either case, they fall outside AST-level manipulation and automata pruning.

Retrieval nondeterminism. When the retrieval source is non-deterministic, the set of candidates explored can differ across runs. Validation catches candidates that fail the current examples, but two runs on the same specification may still produce different regexes that satisfy those examples.

Three directions would strengthen the system. First, the annotated component corpus should be expanded to improve coverage of common categories. Second, larger comparisons against existing regex synthesizers and direct LLM generation would clarify where corpus-guided synthesis adds value. Third, learned ranking from synthesis outcomes could replace the current fixed scoring weights.

8 Conclusion

ReSketch shows that semantic sketches and reusable regex components can make regex synthesis more structured: the user names what each hole should mean, and the system turns that intent into a validated regex through decomposition, search, and example-driven validation. The demonstration on nine tasks shows that the prototype handles both single-hole and multi-hole synthesis under the tested configuration. The decomposition ablation further shows that hole-local evidence is important for the multi-hole tasks in this suite. The prototype is a course-project artifact with clear limitations, but it implements a working path from sketch to validated regex and provides a concrete foundation for future work.

Data Availability

The ReSketch prototype, benchmark tasks, and demonstration scripts are available at <https://github.com/friberk/ReSketch>.

References

- [1] Alberto Bartoli, Giorgio Davanzo, Andrea De Lorenzo, and Eric Medvet. 2014. Automatic Synthesis of Regular Expressions from Examples. *Computer* 47, 12 (2014), 72–80. doi:10.1109/MC.2014.344
- [2] Masudul Hasan Masud Bhuiyan, Berk Cakar, Ethan H. Burmane, James C. Davis, and Cristian-Alexandru Staicu. 2025. SoK: A Literature and Engineering Review of Regular Expression Denial of Service (ReDoS). In *Proceedings of the 20th ACM Asia Conference on Computer and Communications Security*. ACM, New York, NY, USA, 1659–1675. doi:10.1145/3708821.3733912
- [3] Berk Cakar, Charles M. Sale, Sophie Chen, Ethan H. Burmane, Dongyoon Lee, and James C. Davis. 2025. Is Reuse All You Need? A Systematic Comparison of Regular Expression Composition Strategies. doi:10.48550/arxiv.2503.20579
- [4] Qiaochu Chen, Arko Banerjee, Çağatay Demiralp, Greg Durrett, and Isil Dillig. 2023. Data Extraction via Semantic Regular Expression Synthesis. *Proceedings of the ACM on Programming Languages* 7, OOPSLA2 (2023), 1848–1877. doi:10.1145/3622863
- [5] Qiaochu Chen, Xinyu Wang, Xi Ye, Greg Durrett, and Isil Dillig. 2020. Multi-Modal Synthesis of Regular Expressions. In *Proceedings of the 41st ACM SIGPLAN Conference on Programming Language Design and Implementation*. ACM, New York, NY, USA, 487–502. doi:10.1145/3385412.3385988
- [6] James C. Davis, Christy A. Coghlan, Francisco Servant, and Dongyoon Lee. 2018. The Impact of Regular Expression Denial of Service (ReDoS) in Practice: An Empirical Study at the Ecosystem Scale. In *Proceedings of the 2018 26th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*. ACM, New York, NY, USA, 246–256. doi:10.1145/3236024.3236027
- [7] James C. Davis, Louis G. Michael, Christy A. Coghlan, Francisco Servant, and Dongyoon Lee. 2019. Why Aren’t Regular Expressions a Lingua Franca? An Empirical Study on the Re-Use and Portability of Regular Expressions. In *Proceedings of the 2019 27th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*. ACM, New York, NY, USA, 443–454. doi:10.1145/3338906.3338909
- [8] James C. Davis, Francisco Servant, and Dongyoon Lee. 2021. Using Selective Memoization to Defeat Regular Expression Denial of Service (ReDoS). In *2021 IEEE Symposium on Security and Privacy*. IEEE, Piscataway, NJ, USA, 1–17. doi:10.1109/SP40001.2021.00032
- [9] Sumit Gulwani. 2011. Automating String Processing in Spreadsheets Using Input-Output Examples. In *Proceedings of the 38th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. ACM, New York, NY, USA, 317–330. doi:10.1145/1926385.1926423
- [10] Susmit Jha and Sanjit A. Seshia. 2017. A Theory of Formal Synthesis via Inductive Learning. *Acta Informatica* 54, 7 (2017), 693–726. doi:10.1007/s00236-017-0294-5
- [11] Mina Lee, Sunbeom So, and Hakjoo Oh. 2016. Synthesizing Regular Expressions from Examples for Introductory Automata Assignments. In *Proceedings of the 15th International Conference on Generative Programming: Concepts and Experiences*. ACM, New York, NY, USA, 70–80. doi:10.1145/2993236.2993244
- [12] Louis G. Michael, James Donohue, James C. Davis, Dongyoon Lee, and Francisco Servant. 2019. Regexes Are Hard: Decision-Making, Difficulties, and Risks in Programming Regular Expressions. In *2019 34th IEEE/ACM International Conference on Automated Software Engineering*. IEEE, Piscataway, NJ, USA, 415–426. doi:10.1109/ASE.2019.00047
- [13] Kia Rahmani, Mohammad Raza, Sumit Gulwani, Vu Le, David Morris, Arjun Radhakrishna, Gustavo Soares, and Ashish Tiwari. 2021. Multi-Modal Program Inference: A Marriage of Pre-Trained Language Models and Component-Based Synthesis. *Proceedings of the ACM on Programming Languages* 5, OOPSLA, Article 158 (2021), 29 pages. doi:10.1145/3485535
- [14] Armando Solar-Lezama. 2008. *Program Synthesis by Sketching*. Ph.D. Dissertation. University of California, Berkeley.
- [15] Armando Solar-Lezama, Liviu Tancau, Rastislav Bodik, Vijay Saraswat, and Sanjit A. Seshia. 2006. Combinatorial Sketching for Finite Programs. In *Proceedings of the 12th International Conference on Architectural Support for Programming Languages and Operating Systems*. ACM, New York, NY, USA, 404–415. doi:10.1145/1168857.1168907
- [16] Peipei Wang, Chris Brown, Jamie A. Jennings, and Kathryn T. Stolee. 2021. Demystifying Regular Expression Bugs. *Empirical Software Engineering* 27, 1 (2021), 21. doi:10.1007/s10664-021-10033-1
- [17] Xi Ye, Qiaochu Chen, Xinyu Wang, Isil Dillig, and Greg Durrett. 2020. Sketch-Driven Regular Expression Generation from Natural Language and Examples. *Transactions of the Association for Computational Linguistics* 8 (2020), 679–694. doi:10.1162/tacl_a_00339